

## Assignment -2

## DATA VISUALISATION AND PRE-PROCESSING

Assignment Date	28 September 2022
Student Name	GOKULRAJ S
Student Roll Number	612419104011
Maximum Marks	2 Marks

### Question-1:

## Download the Dataset

**SOLUTION:**

[illegible]

### Question-2:

## Loading dataset

**SOLUTION:**

```
df = pd.read_csv("Churn_Modelling.csv")
df
```

```
df = pd.read_csv("Churn_Modelling.csv")
```

[ ] df

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10
...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	9996	15606229	Objiaiku	771	France	Male	39	5	0.00	2	1	0	96270.64
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	1	1	1	101699.77
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0	1	42085.56
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	7507.31	2	1	0	92888.52
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1	1	0	38190.76

10000 rows x 14 columns

### Question-3:

#### 1. Visualizations

##### a) Univariate Analysis

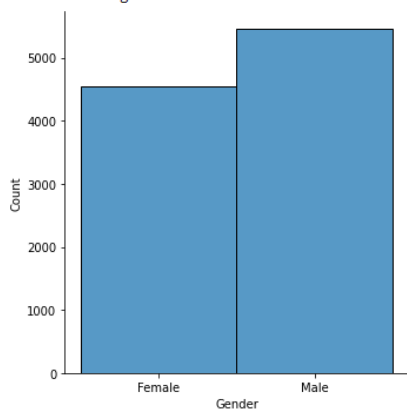
### SOLUTION:

```
sns.displot(df.Gender)
```

##### a) Univariate Analysis

```
[ ] sns.displot(df.Gender)
```

<seaborn.axisgrid.FacetGrid at 0x7f6d4b706350>



##### b) Bi-Variate Analysis

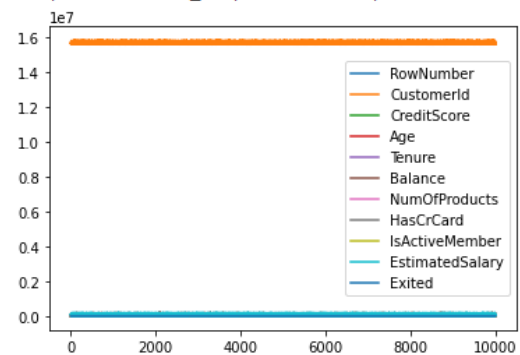
### SOLUTION:

```
df.plot.line()
```

##### b) Bi-Variate Analysis

```
[ ] df.plot.line()
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f6d4b5cf490>

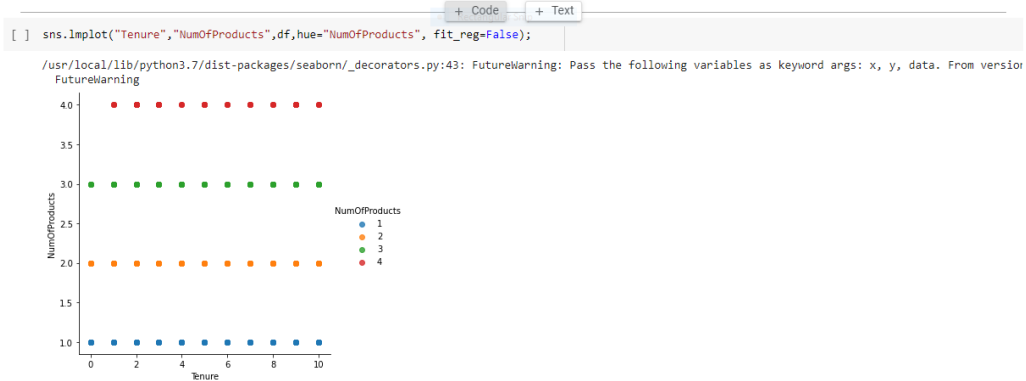


## c) Multi - Variate Analysis

### SOLUTION:

```
sns.lmplot("Tenure", "NumOfProducts", df, hue="NumOfProducts", fit_reg=False);
```

c) Multi - Variate Analysis



## Question-4:

Perform descriptive statistics on the dataset.

### SOLUTION:

```
df.describe()
```

1. Perform descriptive statistics on the dataset.

Rectangular Snip

```
[ ] df.describe()
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.705500	0.515100	100090.239881	
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.581654	0.455840	0.499797	57510.492818	
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.000000	0.000000	11.580000	
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.000000	0.000000	51002.110000	
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.000000	1.000000	100193.915000	
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.000000	1.000000	149388.247500	
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4.000000	1.000000	1.000000	199992.480000	

## Question-5:

Handle the Missing values.

### SOLUTION:

```
data = pd.read_csv("Churn_Modelling.csv")
pd.isnull(data["Gender"])
```

1.Handle the Missing values.

```
[ ] data = pd.read_csv("Churn_Modelling.csv")
pd.isnull(data["Gender"])
```

```
0      False
1      False
2      False
3      False
4      False
...
9995   False
9996   False
9997   False
9998   False
9999   False
Name: Gender, Length: 10000, dtype: bool
```

## Question-6:

Find the outliers and replace the outliers.

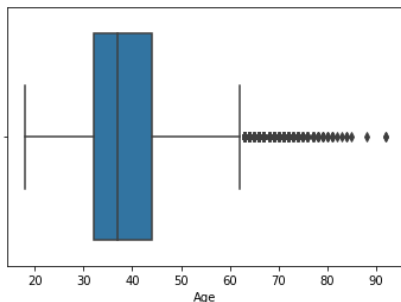
### SOLUTION:

```
sns.boxplot(df['Age'])
```

1.Find the outliers and replace the outliers.

```
[ ] sns.boxplot(df['Age'])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the
FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7f6d4aaf7110>
```



## SOLUTION:

```
df['Age']=np.where(df['Age']>50,40,df['Age'])
df['Age']
```

```
df['Age']=np.where(df['Age']>50,40,df['Age'])
df['Age']
```

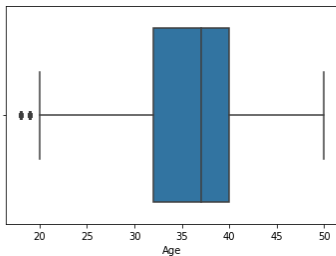
```
0      42
1      41
2      42
3      39
4      43
..
9995   39
9996   35
9997   36
9998   42
9999   28
Name: Age, Length: 10000, dtype: int64
```

## SOLUTION:

```
sns.boxplot(df['Age'])
```

```
[ ] sns.boxplot(df['Age'])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid figure-level keywords are: context, style, and palette.
FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7f6d4b7e8990>
```



## SOLUTION:

```
df['Age']=np.where(df['Age']<20,35,df['Age'])
df['Age']
```

```
[ ] df['Age']=np.where(df['Age']<20,35,df['Age'])
df['Age']
```

```
0      42
1      41
2      42
3      39
4      43
..
9995   39
9996   35
9997   36
9998   42
9999   28
Name: Age, Length: 10000, dtype: int64
```

Question-7:

Check for Categorical columns and perform encoding.

SOLUTION:

```
pd.get_dummies(df, columns=["Gender", "Age"], prefix=["Age", "Gender"]).hea
```

1.Check for Categorical columns and perform encoding.

Rectangular Snip

[ ] pd.get\_dummies(df, columns=["Gender", "Age"], prefix=["Age", "Gender"]).head()

RowNumber	CustomerId	Surname	CreditScore	Geography	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	...	Gender_41	Gender_42	Gende
0	1	15634602	Hargrave	619	France	2	0.00	1	1	1	...	0	1
1	2	15647311	Hill	608	Spain	1	83807.86	1	0	1	...	1	0
2	3	15619304	Onio	502	France	8	159660.80	3	1	0	...	0	1
3	4	15701354	Boni	699	France	1	0.00	2	0	0	...	0	0
4	5	15737888	Mitchell	850	Spain	2	125510.82	1	1	1	...	0	0

5 rows x 45 columns

1.Check for Categorical columns and perform encoding.

Rectangular Snip

pd.get\_dummies(df, columns=["Gender", "Age"], prefix=["Age", "Gender"]).head()

NumOfProducts

HasCrCard	IsActiveMember	...	Gender_41	Gender_42	Gender_43	Gender_44	Gender_45	Gender_46	Gender_47	Gender_48	Gender_49	Gender_50
1	1	1	...	0	1	0	0	0	0	0	0	0
1	0	1	...	1	0	0	0	0	0	0	0	0
3	1	0	...	0	1	0	0	0	0	0	0	0
2	0	0	...	0	0	0	0	0	0	0	0	0
1	1	1	...	0	0	1	0	0	0	0	0	0

### Question-8:

Split the data into dependent and independent variables.

a) Split the data into Independent variables.

### SOLUTION:

```
X = df.iloc[:, :-1].values  
print(X)
```

a) Split the data into Independent variables.

```
[ ] X = df.iloc[:, :-1].values  
    print(X)  
  
[[1 15634602 'Hargrave' ... 1 1 101348.88]  
 [2 15647311 'Hill' ... 0 1 112542.58]  
 [3 15619304 'Onio' ... 1 0 113931.57]  
 ...  
 [9998 15584532 'Liu' ... 0 1 42085.58]  
 [9999 15682355 'Sabbatini' ... 1 0 92888.52]  
 [10000 15628319 'Walker' ... 1 0 38190.78]]
```

b) Split the data into Dependent variables

### SOLUTION:

```
Y = df.iloc[:, -1].values  
print(Y)
```

b) Split the data into Dependent variables.

```
[ ] Y = df.iloc[:, -1].values  
    print(Y)
```

```
[1 0 1 ... 1 1 0]
```

## Question-9:

Scale the independent variables

### SOLUTION:

```
import pandas as pd
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
df[["CustomerId"]] = scaler.fit_transform(df[["CustomerId"]])
print(df)
```

```
import pandas as pd
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
df[["CustomerId"]] = scaler.fit_transform(df[["CustomerId"]])
```

+ Code

+ Text

```
[ ] print(df)
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	\
0	1	0.275616	Hargrave	619	France	Female	42	
1	2	0.326454	Hill	608	Spain	Female	41	
2	3	0.214421	Onio	502	France	Female	42	
3	4	0.542636	Boni	699	France	Female	39	
4	5	0.688778	Mitchell	850	Spain	Female	43	
...	...	...	...	...	...	...	...	
9995	9996	0.162119	Obijiaku	771	France	Male	39	
9996	9997	0.016765	Johnstone	516	France	Male	35	
9997	9998	0.075327	Liu	709	France	Female	36	
9998	9999	0.466637	Sabbatini	772	Germany	Male	42	
9999	10000	0.250483	Walker	792	France	Female	28	

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	
...	...	...	...	...	...	
9995	5	0.00	2	1	0	
9996	10	57369.61	1	1	1	
9997	7	0.00	1	0	1	
9998	3	75075.31	2	1	0	
9999	4	130142.79	1	1	0	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0
...	...	...
9995	96270.64	0
9996	101699.77	0



## Question-10:

Split the data into training and testing

1. List item

### SOLUTION:

```
from sklearn.model_selection import train_test_split

train_size=0.8
X = df.drop(columns = ['Tenure']).copy()
y = df['Tenure']
X_train, X_rem, y_train, y_rem = train_test_split(X,y, train_size=0.8)
test_size = 0.5
X_valid, X_test, y_valid, y_test = train_test_split(X_rem,y_rem, test_size=0.5)
print(X_train.shape), print(y_train.shape)
print(X_valid.shape), print(y_valid.shape)
print(X_test.shape), print(y_test.shape)
```

```
[ ] from sklearn.model_selection import train_test_split
    train_size=0.8
    X = df.drop(columns = ['Tenure']).copy()
    y = df['Tenure']
    X_train, X_rem, y_train, y_rem = train_test_split(X,y, train_size=0.8)
    test_size = 0.5
    X_valid, X_test, y_valid, y_test = train_test_split(X_rem,y_rem, test_size=0.5)
    print(X_train.shape), print(y_train.shape)
    print(X_valid.shape), print(y_valid.shape)
    print(X_test.shape), print(y_test.shape)
```

```
(8000, 13)
(8000,)
(1000, 13)
(1000,)
(1000, 13)
(1000,)
(None, None)
```